

AI Trading Infrastructure – Master Notion Guide

Copy-paste the whole page into Notion. Each top-level heading can become a

Page

Toggle

Database



Overview

- **Mission:** Build a low-latency, contradiction-aware trading infrastructure that unifies *BICEP* (Brownian Inspired Computationally Efficient Parallelization), *Fusion Alpha* (semantic-technical FusionNet pipeline), and the forthcoming *Nimbus Book* (research notebook + paper) into a single, reproducible codebase.
- **Guiding Principles:** Clarity ▶ Modularization ▶ Reproducibility ▶ Robustness ▶ Documentation-first.
- **Quick Links:** README | run_pipeline.py | docs/ | data/raw/.



Roadmap & Milestones

(Use a Notion **Timeline** or **Board** database with the following properties.)

- **Milestone** (title)
- **Due Date** (date)
- **Owner** (person)
- **Status** (select: Idea ▶ In Progress ▶ Blocked ▶ Done)

- **Relevant Module** (multi-select: BICEP, Fusion Alpha, Nimbus Book, Infra)

✨ 2025 Q2 Highlights

- **June 15** — BICEP v1.0 kernel passes unit tests
 - **June 30** — Fusion Alpha walk-forward validation complete
 - **July 10** — Nimbus Book initial chapters drafted
 - **July 31** — Internal alpha of live router on Tesla+8-ticker dataset
-

🔨 Workspaces & Core Databases

1. **Backlog** – Kanban of every task (Task, Description, Module, Priority, Points, Sprint, Status, Dependency).
 2. **Experiments** – Scientific record (Hypothesis, Variant ID, Metrics, Result IMG, Verdict, Merge-candidate?).
 3. **Research Library** – Papers/links with tags (Topic, Relevance, Read By, Notes).
 4. **Decision Log** – Architecture & trade-off log (Date, Decision, Context, Alternatives, Consequences).
-

🧠 Module: BICEP

- **Goal:** Ultra-efficient, Brownian-inspired data-parallel kernels for PyTorch & CUDA.
- **Key Files:** src/bicep/kernels/*.cu, train_bicep.py, tests/test_bicep.py.
- **Implementation Checklist:**
 - Finalise random walk bit-mapping logic
 - GPU parallelisation benchmarks (RTX 4090 vs A100)
 - Integrate into FusionNet as drop-in replacement
 - Document theoretical derivation (Nimbus Book §2)
- **Validation Metrics:** TFLOPS utilisation, epoch time, loss parity with baseline.

Module: Fusion Alpha

- **Goal:** Fuse FinBERT sentiment embeddings with technical indicators to detect price-sentiment contradictions and generate alpha.
- **Key Scripts:** prepare_dataset.py, train_fusion.py, tune_fusion.py, predict_fusion.py.
- **Pipeline Stages:**
 1. **Data Ingest** – data/raw/*.csv → OHLCV + news scrape
 2. **Feature Build** – technical_features, finbert_embeddings
 3. **FusionNet Training** – Multi-head attention + contradiction severity
 4. **Evaluation** – k-fold CV, walk-forward, benchmark strategies
 5. **Live Router** – live_router.py auto-fetch + inference
- **KPIs:** Sharpe, max drawdown, hit-rate vs benchmark.

Module: Nimbus Book

- **Form:** Jupyter-Book in docs/, exported as static site & PDF.
- **Sections:**
 - Intro & Literature Survey
 - BICEP Theoretical Foundations
 - Fusion Alpha Methodology
 - Experiments & Ablation Studies
 - Future Work & PhD Directions

Development Workflow

1. Environment Setup

```
conda env create -f environment.yml # or pip install -r requirements.txt
pre-commit install                  # enforce hooks
```

- 1.
 2. **Branching Model** – main | dev | feat/* | fix/* | exp/* (no direct commits to main).
 3. **Commit Messages** – Conventional Commits (feat(bicep): add cuda kernel #24).
 4. **CI/CD** – GitHub Actions: lint → test → build Docker → deploy docs.
 5. **Code Review** – 1 approving review + green CI required.
-

Data Management

- **Versioning:** DVC with remote on S3 (dvc remote add -d s3 s3://ai-trading-data).
 - **Layers:** raw → interim → processed → features → predictions.
 - **Schemas:** store as Parquet w/ schema files in schemas/.
 - **Privacy:** Strip PII, encrypt sensitive feeds with age-encryption.
-

Folder Structure Explained

```
.
├── data/           # Datasets & embeddings
│   ├── raw/        # Unmodified source CSVs
│   ├── interim/    # Cleaned but not feature-engineered
│   ├── processed/  # Analysis-ready
│   └── features/   # Numpy/torch arrays
├── models/        # Saved .pth checkpoints & model cards
│   ├── checkpoints/
│   └── experiments/
├── src/           # Library code
│   ├── bicep/
│   ├── fusion_alpha/
│   ├── nimbus_book/
│   └── common/
```

```
|— notebooks/      # Exploratory analysis
|— scripts/        # CLI helpers (e.g., run_pipeline.sh)
|— tests/          # pytest suites
|— docker/         # Dockerfiles & compose
|— docs/           # Auto-generated site (Nimbus Book)
```

Usage Ritual

1. git pull --rebase
2. Activate env → dvc pull → python scripts/run_pipeline.py --stage ingest
3. Develop feature → pytest -q → pre-commit run --all-files.
4. Push PR → code review → merge.
5. Trigger deployment (docs + Docker image tag).

CLI Reference

| Tip:

- **Run full pipeline**

```
python run_pipeline.py --ticker TSLA --start 2015-01-01 --mode train
```

- **Hyperparameter tuning**

```
python tune_fusion.py --trials 50 --sampler tpe
```

- **Live route**

```
python live_router.py --tickers TSLA AAPL NVDA --interval 1m --log live.log
```

- **Benchmark**

```
python evaluate_strategy.py --baseline buy_and_hold
```



Daily Logs & Meeting Notes Template

Daily Stand-Up (YYYY-MM-DD)

- **Yesterday:**
- **Today:**
- **Blockers:**
- **Next Milestone:**

Glossary

| Term | Meaning |
|---------------------------------|---|
| BICEP | Brownian Inspired Computationally Efficient Parallelisation |
| Fusion Alpha | Sentiment-Technical fusion network for alpha generation |
| Nimbus Book | Comprehensive research notebook/documentation |
| Contradiction Score (CS) | Measure of sentiment-price divergence |

Quick-Access Checklists

- Conda env recreated after any environment.yml change
- dvc push after adding data artefacts
- Draft decision logged within 24 h of change
- PR merged \leq 250 LOC unless justified

Future Enhancements

- Integration with FPGA low-latency order gateway
- Real-time FinBERT inference via ONNX-Runtime
- Switch to duckDB for local columnar analytics
- Deploy on k8s cluster with GPU node autoscaling

| End of Guide – happy building!

